

DEEPMINDQ CRM · MS6 DESIGN FOUNDATION

Complete **Supporting** Documentation

Five comprehensive design documents establishing the UX philosophy, information architecture, component specifications, interaction rules, and emotional copy library for the DeepMindQ CRM platform. Locked foundation for MS7 through MS11 development.

UX Philosophy · Information Architecture · Component Specifications
Interaction Rules · Emotional Copy Library

Version 1.0 · MS6 Stage 1 · Confidential

DEEPMINDQ CRM · INTERNAL DESIGN DOCUMENTATION

Table of Contents

Table of Contents	1
Chapter 1: UX Philosophy Document	3
1.1 The DeepMindQ Experience Promise	3
1.2 Design North Star: The VP Sales Mental Model	3
1.3 Core Philosophy: Intelligence as Executive Briefing	4
1.4 The 5-Question Morning Intelligence Protocol	4
1.5 The 15 Locked Core Design Patterns	5
Chapter 2: Information Architecture Documentation	7
2.1 Navigation Model: Intelligence Hub First	7
2.2 The Seven Core Experiences	7
2.3 Progressive Disclosure Model (L1–L4)	8
2.4 Navigation Transformation: Current to Future	8
Chapter 3: Component Specifications	10
3.1 ConfidenceIndicator — INCORPORATE	10
3.2 EvidenceChain — INCORPORATE	10
3.3 InlineReasoning → INCORPORATE	11
3.4 IntelligenceNarrative — REDESIGN	11
3.5 ProgressiveDisclosure — REDESIGN	11
3.6 RecommendationCard — REDESIGN	12
3.7 Unified Component Architecture	12

Chapter 4: Interaction Rules 13

4.1 Progressive Disclosure Interaction Rules	13
4.2 Navigation Interaction Rules	13
4.3 Trust Visualization Interaction Rules	13
4.4 Error and Empty State Interaction Rules	14
4.5 State Transition Model	14

Chapter 5: Emotional Copy Library 16

5.1 Voice and Tone Principles	16
5.2 Signal and Intelligence Copy Patterns	16
5.3 Error Message Copy Patterns	16
5.4 Empty State Copy Patterns	17
5.5 Success and Milestone Copy Patterns	17
5.6 Action and CTA Copy Patterns	17

Chapter 1: UX Philosophy Document

This document establishes the foundational UX philosophy for the DeepMindQ CRM platform, as defined and locked during MS6 Design Foundation. Every design decision, component specification, interaction pattern, and visual language choice in MS7 through MS11 must trace back to the principles articulated here. This is not a style guide or a component catalog; it is the philosophical framework that gives those artifacts their meaning and coherence.

1.1 The DeepMindQ Experience Promise

The DeepMindQ Experience Promise is the single most important artifact in this document. It is the contract between the product and its users. Every feature, every interaction, every visual element either serves this promise or does not belong in the product. The promise is not a feature list or a technical specification. It is an outcome guarantee.

"Every morning, a VP Sales opens DeepMindQ and within 3 minutes knows exactly what changed, why it matters, who to engage, what to say, and what to do — with full confidence in the evidence behind every recommendation."

This promise contains three measurable commitments. First, the time commitment: 3 minutes to complete situational awareness. This is not an aspirational target; it is a design constraint. Every element of the Intelligence Hub must be architected to deliver its value within this window. Second, the completeness commitment: the user must have answers to all five morning questions (What changed? Why does it matter? Who should we engage? What should we say? What should we do?) without navigating away from the primary experience. Third, the trust commitment: every recommendation must carry its evidence footprint. The user should never have to guess why the system suggests something; the reasoning, sources, and confidence levels must be transparently available.

The three pillars that support this promise are Intelligence, Not Data; Confidence at Every Layer; and the 3-Morning-Minute Protocol. Intelligence, Not Data means the system delivers intelligence, not information. Every element answers what, why, who, how, and why trust this. Confidence at Every Layer means every recommendation carries an evidence footprint, and users never have to guess why the system suggests something. The 3-Morning-Minute Protocol means the entire Intelligence Hub experience is designed to deliver complete situational awareness in under 3 minutes.

1.2 Design North Star: The VP Sales Mental Model

Every UX decision in DeepMindQ passes through one filter: Would a VP Sales trust this, rely on this daily, and recommend this to their team? The VP Sales is not a demographic or a persona in the traditional UX sense. It is a decision-making mindset that represents our primary user archetype. This persona embodies three behavioral principles that shape the entire experience.

The first principle is the Time-Scarce Decision Maker. A VP Sales arrives each morning with approximately 15 minutes maximum for intelligence review. They do not browse dashboards, they do not read feeds, and they do not explore interfaces. They need signal, not noise. Every element of the DeepMindQ experience must justify its cognitive cost. If a component requires more than 5 seconds of processing time, it must deliver proportional value. Progressive disclosure exists precisely to manage this: the morning briefing gives the summary; deeper layers are available for those who need them.

The second principle is the Trust-Sensitive Executive. This persona will not act on AI recommendations without understanding why. Evidence is not optional for this user; it is a prerequisite for action. The trust-sensitive executive has been burned by AI tools that make confident-sounding but incorrect predictions. DeepMindQ must build trust through radical transparency: every confidence score, every evidence chain, every reasoning step must be immediately accessible. The system earns trust by showing its work, not by asking for it.

The third principle is the Action-Oriented Leader. Intelligence without a recommended next step is wasted cognitive load for this persona. Every insight in DeepMindQ must terminate in a clear, actionable recommendation. This does not mean the system tells the user what to do; rather, it suggests what to consider, with enough context for the user to make an informed decision. The distinction between direction and briefing is critical here and is explored in depth in the next section.

1.3 Core Philosophy: Intelligence as Executive Briefing

This is Design Pattern #1 and it overrides all other patterns. DeepMindQ does not tell the user what to do. DeepMindQ briefs the user on what is happening, why it matters, and recommends what to consider — with full evidence transparency. This distinction is subtle but critical. An executive briefing gives the reader context, analysis, and recommendations, but the executive makes the decision. This is the relationship DeepMindQ must have with its users.

The philosophy manifests in four key distinctions. Briefing, Not Direction means the system presents intelligence as a structured, prioritized, evidence-backed briefing document. The decision always belongs to the human. Recommendation, Not Instruction means AI suggests actions with confidence levels and evidence trails. Users evaluate, adjust, and decide. The system provides the briefing; the user provides the judgment. Evidence, Not Authority means every insight carries its evidence footprint. Trust is built through transparency, not through brand or authority. Context, Not Complexity means progressive disclosure reveals depth on demand. The morning briefing gives the summary; deeper layers are available for those who need them. Intelligence surfaces first; complexity reveals itself when needed.

1.4 The 5-Question Morning Intelligence Protocol

The 5-Question Framework is the structural backbone of the Intelligence Hub. It is not a feature; it is an information architecture constraint. Every card, signal, recommendation, and action in the Hub must map to at least one of these five questions. If a component cannot answer which question it serves, it belongs in a deeper experience layer, not the Hub. This constraint prevents the Hub from becoming an overloaded dashboard.

Question	maps to	Information Type
1. What Changed?	Surveillance	Critical signals, account movements, market shifts, competitive changes, ranked by impact
2. Why It Matters?	Analysis	Context, reasoning, impact analysis, trend correlation, revenue implication
3. Who to Engage?	Targeting	Account-specific recommendations, priority-ranked engagement targets with reasoning
4. What to Say?	Preparation	Conversation angles, talking points, contextual intelligence for each engagement
5. What to Do?	Action	Recommended actions with confidence scores, timelines, evidence trails

The design rule is absolute: No intelligence element exists in the Hub that does not map to at least one of these five questions. This constraint ensures the Intelligence Hub remains focused, actionable, and bounded. The 75+ deeper screens that exist in the current system remain accessible through progressive disclosure within the seven core experiences, but they are not surfaced in the Hub unless they directly answer one of the five questions.

1.5 The 15 Locked Core Design Patterns

MS6 locks 15 core design patterns that form the DNA of the DeepMindQ experience. These patterns are selected from the comprehensive 72-point behavioral psychology checklist and represent the highest-impact design interventions. Each pattern is categorized by its primary domain: Philosophy, Architecture, Trust, Emotional, or Visual. Every component, interaction, and visual element in MS7 through MS11 must map to at least one of these patterns.

#	Pattern	Category	Core Principle
1	AI Intelligence as Executive Briefing	Philosophy	AI briefs; humans decide
2	Cognitive Load Reduction	Philosophy	Minimize mental effort per decision
3	Progressive Disclosure	Architecture	Summary to depth on demand
4	Navigation Intelligence	Architecture	Nav answers what to focus on next
5	Decision Fatigue Prevention	Philosophy	Reduce choices per decision point
6	Trust Visualization / Evidence Footprint	Trust	Every insight shows its proof
7	Confidence + Freshness + Source Transparency	Trust	Confidence is quantified and sourced
8	Emotional Copywriting	Emotional	Words carry emotional weight

#	Pattern	Category	Core Principle
9	Compassionate Error Handling	Emotional	Errors explain why, offer a path forward
10	Loss Aversion Design	Emotional	Frame as risk prevention
11	Endowment Effect / Personal Ownership	Emotional	Users feel ownership of their workspace
12	Premium Enterprise Visual Language	Visual	Dark glass-morphism, calm aesthetics
13	Human Assistance Layer	Architecture	AI + human collaboration model
14	Empty State + Success State Experience	Visual	Every state has meaningful content
15	Component Architecture System	Architecture	Atomic design: tokens, atoms, molecules, organisms


The remaining patterns from the 72-point checklist will be applied progressively through MS7–MS11 as specific screens are developed. Each screen development must pass a 5-question gate: Is this important? What should the user do? Why trust this? Does the user have control? Does this match the Design DNA?

Chapter 2: Information Architecture Documentation

This document specifies the information architecture for DeepMindQ as established in MS6. It covers the navigation model, the Intelligence Hub structure, the progressive disclosure system, and the relationship between the seven core experiences and the 75+ deeper screens. This IA is locked and serves as the structural foundation for all MS7–MS11 screen development.

2.1 Navigation Model: Intelligence Hub First

The Intelligence Hub is the default landing page and the primary experience. It serves as the Executive Briefing Room. Every session starts at the Intelligence Hub. No other screen is the default entry point. The Hub is not a "one screen to rule them all" overloaded dashboard. It is a focused briefing environment that answers the 5 morning questions within a 3-minute protocol. Other experiences exist as deeper intelligence environments accessible from the Hub.



The architectural principle is: Surface intelligence first. Reveal complexity only when needed. The 75+ deeper screens remain as progressive disclosure depth inside these seven experiences. They should NOT become top-level navigation items. Navigation shows 7 experiences; depth lives inside each experience. This reduces cognitive load from 76 choices to 7, with depth available on demand. The navigation itself is intelligent (Design Pattern #4): it answers the question "What should I focus on next?" through contextual badges, priority indicators, and attention guidance.

2.2 The Seven Core Experiences

Each experience is a self-contained intelligence environment with its own information hierarchy, interaction patterns, and progressive disclosure depth. The following table documents each experience, its primary purpose, the morning questions it answers, and its key components.

Experience	Primary Purpose	Morning Questions	Key Components
Intelligence Hub	Executive briefing room; answers all 5 questions	Priority Signals, Recommended Actions, Evidence Summaries	

Experience	Primary Purpose	Morning Questions	Key Components
Market Intelligence	Market signals, industry trends	1. What's new in the market?	Signal Hub, Trend Analysis, Competitive Tracker, Market Segments
Account Intelligence	Account-specific deep intelligence	2. How is our key account performing?	Account Enterprise, Opportunity Radar, Contact Intelligence, Engagement Metrics
AI Advisors	AI-powered advisory conversations	3. What are the risks and opportunities?	AI Insights, Advisory Sessions, Scenario Analysis, Recommendations
Workspace	Personal workspace for saved views	4. What are my key priorities?	Saved Views, Custom Briefings, Personal Preferences, Workspace Settings
Operations	System health, data management	5. How is the system performing?	System Health, Data Sources, Import Management, Audit Logs
Command Center	Operational command view	6. What's the overall intelligence picture?	Intelligence Queue, Action Pipeline, Status Metrics, Operational Alerts

2.3 Progressive Disclosure Model (L1–L4)

Progressive disclosure is Design Pattern #3 and the mechanism that makes Intelligence as Executive Briefing work. Not every user needs the same depth. A VP Sales who trusts the system after weeks of use may never drill past Level 1. A new user or a skeptic will want Level 3 evidence. The key insight: each level is independently accessible and answers a specific cognitive need. The experience adapts to the user's trust level.

Level	Name	Cognitive Question	Contents	Target User
L1	Decision Layer	Do I need to act on this?	Headline, confidence score, Priority badge, (Invest, Hold, Sell)	Executives, Analysts, Investors
L2	Reasoning Layer	Why does the system think this?	Reasoning, positive/negative evaluation factors	Analysts, Researchers, Strategists
L3	Evidence Layer	Show me the proof	Evidence chain with sources, Skepticism, validation links	Analysts, Researchers, Compliance
L4	Exploration Layer	What else should I know?	Related signals, connected insights, adjacent opportunities	Researchers, Analysts, Strategists

Each level is independently accessible. Users who trust the system stay at L1. Skeptics drill to L3. Researchers explore L4. The experience adapts to the user's trust level — the user is never forced to digest complexity they do not need.

2.4 Navigation Transformation: Current to Future

The current navigation structure presents 16 items across 5 sections in the sidebar, with 80+ screens accessible through various navigation paths. The routing mechanism uses hash-based SPA navigation via Zustand store (‘activeView: ViewId’). This creates significant cognitive load: a new user faces 16 navigation choices without clear prioritization, and the relationship between screens is not immediately apparent.

The future navigation model collapses 16 items into 7 experiences, with the Intelligence Hub as the default landing page. The 75+ deeper screens remain accessible through progressive disclosure within each experience, not as navigation items. The routing mechanism will be updated to use a hierarchical navigation model where each experience can contain sub-views, but the primary navigation surface shows only 7 items. This transformation reduces the initial cognitive load from 16 choices to 7, with depth available on demand within each experience. The navigation itself becomes intelligent, showing contextual indicators of what needs attention across all experiences.

Dimension	Current State	Future State (MS6+)
Navigation items	16 sidebar items across 5 sections	7 experiences, Hub as default
Screen access	80+ screens via flat navigation	75+ screens via progressive disclosure
Default landing	Intelligence Operations Center	Intelligence Hub (Executive Briefing)
Cognitive load	16 choices at first interaction	7 choices at first interaction
Navigation intelligence	Static sidebar with section labels	Contextual priority indicators
Depth access	All screens equal priority	Progressive disclosure L1–L4

Chapter 3: Component Specifications

This document evaluates the six existing Intelligence OS components against MS6 design principles. Each component is assessed for alignment with the core patterns: cognitive load reduction, intelligence as executive briefing, trust visualization, premium enterprise experience, and human-centered interaction. Components that align are incorporated into the Design System. Components that conflict are redesigned. Components are never blindly preserved or discarded; every decision is traceable to a design principle.

3.1 ConfidenceIndicator — INCORPORATE

File: src/components/intelligence-os/confidence-indicator.tsx (184 lines)

Verdict: Evaluate → Incorporate. Strong alignment with MS6 principles.

Action: Incorporate into Design System as universal confidence display.

The ConfidenceIndicator is the most aligned of all six components. It provides 4 visual modes (ring, bar, badge, score) and 4 size variants (xs, sm, md, lg), with animated value transitions and confidence tier colors from the existing design token system. The component directly serves Design Pattern #7 (Confidence + Freshness + Source Transparency) and Design Pattern #6 (Trust Visualization / Evidence Footprint).

Current interface: ConfidenceIndicatorProps with value (0-100), mode (ring|bar|badge|score), label, showPercentage, size (xs|sm|md|lg), animated, className. The component uses the existing design token system for confidence tier colors (high ≥ 70 = green, medium ≥ 45 = amber, low = red) and includes animated transitions. No changes to the core interface are needed. Visual styling will be updated to match the premium enterprise visual language, but the interaction model and prop interface remain unchanged.

3.2 EvidenceChain — INCORPORATE

File: src/components/intelligence-os/evidence-chain.tsx (180 lines)

Verdict: Evaluate → Incorporate. Core trust pattern.

Action: Redesign visual styling to match premium enterprise language.

The EvidenceChain is a core trust visualization component that displays a numbered evidence trail with source icons, date stamps, relevance scores, and verification links. It serves Design Pattern #6 (Trust Visualization / Evidence Footprint) directly. The component uses the existing design token system for theming and supports compact and full display modes.

Current interface: EvidenceChainProps with items (EvidenceChainItem[]), title, conclusion, verdict (strong|moderate|weak), compact, className. Each item has source, sourceType (news|filing|web|database|social|internal|sec|press), snippet, url, date, relevanceScore. The visual design will be updated to match the premium enterprise visual language (glass-morphism cards, semantic color coding, refined typography), but the interaction model and data interface remain unchanged.

3.3 InlineReasoning → INCORPORATE

File: src/components/intelligence-os/inline-reasoning.tsx (158 lines)

Verdict: Evaluate → Incorporate. Key L2 reasoning surface.

Action: Refine factor visualization, add expand-to-evidence link pattern.

The InlineReasoning component is the unified L2 reasoning surface that shows "Why does the system think this?" directly within narratives. It displays a primary reasoning statement, positive and negative confidence factor tags (green trending-up for positive, amber alert-triangle for negative), and an optional "See full evidence" expand link. It serves Design Pattern #3 (Progressive Disclosure) by providing the reasoning layer between the decision layer and the evidence layer.

Current interface: InlineReasoningProps with reasoning, positiveFactors (string[]), negativeFactors (string[]), onClickExpand, compact, showExpandLink. The component will be refined with improved factor tag visualization (larger hit targets, clearer positive/negative distinction), better transition animations for the expand link, and integration with the new evidence chain component at L3.

3.4 IntelligenceNarrative — REDESIGN

File: src/components/intelligence-os/intelligence-narrative.tsx (811 lines)

Verdict: Evaluate → Redesign. Right pattern, wrong complexity.

Action: Decompose into composable atoms, reduce API surface.

The IntelligenceNarrative is the most complex component at 811 lines. It is the core intelligence delivery vehicle — a narrative-first, action-terminated, progressive disclosure surface supporting 6 variants (signal, opportunity, risk, enrichment, reasoning, action) and L1–L4 progressive disclosure. The concept is perfectly aligned with Design Patterns #1, #3, and #6. However, the implementation violates Design Pattern #2 (Cognitive Load Reduction): the component is a monolith that handles too many concerns.

The redesign plan decomposes IntelligenceNarrative into composable atoms: NarrativeHeader (L1: headline, confidence, badge, timestamp), NarrativeReasoning (L2: reasoning text, factor tags), NarrativeEvidence (L3: evidence chain), NarrativeExploration (L4: related signals, actions), and NarrativeAction (action CTA footer). The composite IntelligenceNarrative assembles these atoms based on the variant and disclosure level. This decomposition reduces the component from 811 lines to approximately 5 atoms of 80–120 lines each, plus a 120-line composite. The API surface shrinks from ~25 props to ~8 essential props, with variant-specific props handled internally by the composite.

3.5 ProgressiveDisclosure — REDESIGN

File: src/components/intelligence-os/progressive-disclosure.tsx (361 lines)

Verdict: Evaluate → Redesign. Concept aligns, implementation inconsistent.

Action: Unify with token system, merge with IntelligenceNarrative pattern.

The ProgressiveDisclosure component provides a 4-level progressive disclosure pattern (L1–L4) with tab-like layer toggles and Framer Motion animations. The concept is perfectly aligned with Design Pattern #3. However, the implementation uses inline Tailwind colors (hardcoded hex values like #059669, #f59e0b, #ef4444) instead of the design token system. This creates a consistency gap with the rest of the Intelligence OS components, which use design-tokens.ts for all theming.

The redesign plan has two parts. First, migrate all hardcoded colors to the design token system (semantic confidence colors, domain variant colors, surface elevations). Second, evaluate whether ProgressiveDisclosure should remain a standalone component or merge into the decomposed IntelligenceNarrative pattern. Given that both components implement the same L1–L4 disclosure model, the preferred approach is to make ProgressiveDisclosure use the same atomic components (NarrativeHeader, NarrativeReasoning, NarrativeEvidence, NarrativeExploration) that the redesigned IntelligenceNarrative uses, eliminating the parallel implementation.

3.6 RecommendationCard — REDESIGN

File: src/components/intelligence-os/recommendation-card.tsx (524 lines)

Verdict: Evaluate → Redesign. Too monolithic, violates separation of concerns.

Action: Split into data layer and presentation layer.

The RecommendationCard displays AI-generated account recommendations with a comprehensive progressive disclosure structure: header, "Why This Account" section, expandable evidence list, risk factors, recommended action box, confidence breakdown, enrichment indicators, and stats footer. The concept serves Design Patterns #1, #3, #6, and #7. However, at 524 lines, the component violates separation of concerns by fetching its own data from `/api/recommendations/${companyId}` on mount.

The redesign plan separates the component into two layers. The data layer (RecommendationData) is responsible for fetching, normalizing, and caching recommendation data. The presentation layer (RecommendationCard) is a pure presentation component that receives normalized data as props and renders using the atomic component library (NarrativeHeader, ConfidenceIndicator, EvidenceChain, InlineReasoning). This separation enables better testing, independent evolution of data fetching and presentation, and reusability of the presentation layer with different data sources.

3.7 Unified Component Architecture

After evaluating and redesigning all six components, the unified component architecture follows atomic design principles: Design Tokens (single source of truth), Intelligence Atoms (ConfidenceIndicator, EvidenceChain, InlineReasoning, ActionCTA, StatusBadge), Intelligence Molecules (IntelligenceNarrative, RecommendationCard, HeroNarrative, IntelligenceBriefing), and Intelligence Organisms (IntelligenceHub, Market Intelligence, Account Intelligence, Command Center). This architecture ensures consistency, composability, and testability across all seven experiences.

Chapter 4: Interaction Rules

This document defines the interaction rules for the DeepMindQ experience. These rules govern how components behave, how state transitions work, how users interact with intelligence elements, and how the system responds to user actions. Every interaction in MS7–MS11 must conform to these rules. They are derived from the 15 core design patterns and the behavioral psychology checklist.

4.1 Progressive Disclosure Interaction Rules

Progressive disclosure is the primary interaction mechanism. It governs how intelligence is revealed across all four layers. The following rules define the behavior of every disclosure-enabled component in the system.

Rule ID	Rule	Behavior
PD-01	Independent Level Access	Each L1–L4 level is independently accessible. Users can jump directly to any level with a single click.
PD-02	State Persistence	Disclosure state persists within a session. If a user expands L2 on a signal, L2 remains expanded until the user collapses it or the session ends.
PD-03	Animation Consistency	All disclosure transitions use Framer Motion AnimatePresence with 200ms duration and easeOut.
PD-04	Cognitive Load Gate	No more than 3 expanded disclosure levels visible simultaneously. Expanding a new level collapses the least recently used expanded level.
PD-05	Trust Adaptation	The system tracks which levels a user accesses most frequently and pre-optimizes the disclosure state for the next session.

4.2 Navigation Interaction Rules

Navigation serves Design Pattern #4 (Navigation Intelligence). The navigation does not just route users between views; it answers the question "What should I focus on next?" through contextual indicators and attention guidance.

Rule ID	Rule	Behavior
NAV-01	Hub as Default	Every new session starts at the Intelligence Hub. No other experience is the default entry point.
NAV-02	Attention Indicators	Navigation items show contextual badges when they contain unread intelligence: count, priority, or status.
NAV-03	Depth Breadcrumb	When a user drills into a deeper screen within an experience, a breadcrumb trail shows the path back to the parent view.
NAV-04	Session Context	Navigation preserves session context. Switching between experiences does not reset the current view or filters.
NAV-05	Keyboard Navigation	Full keyboard navigation: Tab between experiences, Enter to select, Escape to return to the previous view.

4.3 Trust Visualization Interaction Rules

Trust visualization governs how confidence, evidence, and reasoning are displayed and interacted with. These rules ensure that trust-building interactions are consistent, transparent, and accessible across all experiences.

Rule ID	Rule	Behavior
TR-01	Confidence Always Visible	Every intelligence element displays its confidence score at L1. Confidence is never hidden.
TR-02	Evidence on Demand	Evidence chains are accessible at L3 via progressive disclosure. Hovering on a confidence score reveals the evidence chain.
TR-03	Source Verification	Every evidence item includes a "Verify" link that opens the original source in a new tab.
TR-04	Reasoning Transparency	Reasoning is accessible via inline expansion. Positive and negative confidence factors are clearly labeled.
TR-05	Freshness Indication	Every intelligence element shows its data freshness: "Updated 2 hours ago" or "Data from 2023-10-27".

4.4 Error and Empty State Interaction Rules

Error and empty states are critical trust moments. Design Pattern #9 (Compassionate Error Handling) and Design Pattern #14 (Empty State + Success State Experience) govern these interactions. The system must never leave the user in a dead end.

Rule ID	Rule	Behavior
ES-01	No Dead Ends	Every error state includes a path forward: "Retry", "Use cached data", "Configure data sources".
ES-02	Human Error Messages	Error messages explain why in human terms: "We could not analyze this account because it is private." vs "Invalid API key".
ES-03	Empty States Inspire	Empty states describe what the user could do: "No signals detected yet. Connect a data source to get started."
ES-04	Success States Celebrate	Success states acknowledge meaningful outcomes: "3 new signals detected. 2 high-priority alerts." vs "Data received successfully."
ES-05	Loading States Inform	Loading states indicate what is being loaded and approximate time: "Analyzing 12 accounts (2 min)." vs "Loading..."

4.5 State Transition Model

Every intelligence element in the system follows a consistent state transition model. Understanding these transitions is essential for implementing smooth, predictable interactions across all experiences. The model defines six primary states and the transitions between them.

State	Trigger	Visual Indicators	User Actions Available
Loading	Component mount, data fetch initiated	Skeletonized content, pulse animation, spinner	Cancel, retry, timeout
Ready	Data received, intelligence feed populated	Full display with confidence scores	Expand L2, dismiss, action
Expanded L2	User clicks reasoning/expand	Reasoning panel with factors	Collapse L2, expand L3
Expanded L3	User clicks evidence/expand	Evidence chain with sources	Verify source, expand L4, collapse
Expanded L4	User clicks explore/expand	Related signals and actions	View L4, navigate to related, collapse

State	Trigger	Visual Indicators	User Actions Available
Error	Data fetch failure, computation error	Error icon with human message and icon	Retry, undo step, reconfigure, contact support

Chapter 5: Emotional Copy Library

This document provides the initial Emotional Copy Library for DeepMindQ. It establishes the tone, voice, and specific copy patterns for all user-facing text in the product. Every word in the interface carries emotional weight (Design Pattern #8). Error messages acknowledge frustration before offering solutions. Success states celebrate outcomes, not features. Empty states inspire possibility, not abandonment. The copy speaks to the human, not the feature.

5.1 Voice and Tone Principles

The DeepMindQ voice is calm, confident, and precise. It speaks to an executive audience that values clarity over personality, substance over style, and trust over charm. The voice avoids hype, exaggeration, and unnecessary enthusiasm. It does not say "Amazing insights detected!" or "You won't believe what we found!" It says "3 high-priority signals detected" and "Here is what changed and why it matters." The tone adapts to context: neutral for standard intelligence, urgent for critical signals, supportive for errors, and encouraging for empty states.

The four voice attributes are: Authoritative but not arrogant (the system presents intelligence with evidence, not authority), Calm but not cold (the system cares about the user's success, expressed through clarity and helpfulness, not through enthusiasm or emotion), Precise but not pedantic (the system uses specific language that conveys exact meaning without unnecessary jargon or over-explanation), and Supportive but not patronizing (the system helps without telling the user what to do, respects the user's intelligence, and assumes competence).

5.2 Signal and Intelligence Copy Patterns

Signals and intelligence elements use structured copy patterns that communicate urgency, impact, and actionability. The following patterns are locked for MS7–MS11 development.

Context	Copy Pattern	Example
Critical signal detected	[Entity] [change] [impact] [Recommendation]	Account relationship change detected. This may affect your renewal conversation.
High-priority opportunity	[Entity] shows [opportunity signal] [Evidence] [Confidence]	Market signals evidence buying intent. 3 recent job postings for your company.
Risk alert	[Entity] [risk indicator] [Potential impact] [Mitigation suggestion]	Operational impact. Migration suggested. Historical patterns suggest 60-day inactivity.
Routine intelligence	[Entity] [update]. [Significance]	Data source updated their pricing page. Your opportunity score increased by 5%.

5.3 Error Message Copy Patterns

Error messages follow the compassionate error handling pattern (Design Pattern #9). Every error message has three parts: acknowledge the situation, explain why in human terms, and offer a clear next step. Error messages never use technical jargon as the primary text, never display stack traces or error codes to the user,

and never leave the user without a path forward.

Error Type	Acknowledgment	Explain	Next Step
Data fetch failure	We could not load intelligence for this account.	Intelligence for this account is not updated every 14 days. The data source is configured to update every 14 days.	Refresh the data source or contact your administrator.
AI computation error	We could not generate a recommendation.	AI recommendation generation failed. This may be due to a temporary issue with the AI engine or a configuration error.	Wait a few minutes and try again. If the error persists, contact your administrator.
Network timeout	The request took longer than expected to complete.	Our servers are experiencing a delay in processing your request.	Retry the request. If the delay persists, contact your administrator.
Permission denied	You do not have access to this intelligence.	This intelligence is restricted to administrators only.	Contact your administrator to request access.

5.4 Empty State Copy Patterns

Empty states describe what the user could do, not what is absent. They inspire possibility and provide clear paths to value. Empty states never use generic "No data" messages and never make the user feel like they have done something wrong.

Context	Copy Pattern	Example
No signals	No [type] detected yet.	No market signals detected yet. Connect a data source to start receiving real-time market data.
No recommendations	No recommendations available.	No recommendations available for this account. We need more engagement data to generate personalized recommendations.
No search results	No results for "[query]".	No results for "enterprise AI". Try broader terms like "AI" or "machine learning".
No saved views	You have not saved any views yet.	You have not saved any views yet. Save your current view to quickly access your favorite data later.

5.5 Success and Milestone Copy Patterns

Success states and milestone recognitions follow the enterprise celebration style: dignified, measured, and meaningful. The system acknowledges achievements with premium milestone recognition, never with confetti, sparkle, or flashy animations. Milestones are recognized because they represent genuine business outcomes, not because the user clicked a button.

Milestone Type	Recognition Copy	Visual Treatment
First intelligence-driven deal	Your first intelligence-driven deal.	Dignified notification. The signal that the first deal was closed.
Streak: 7 consecutive days	7 days of consistent intelligence gathering.	Subtle counterstreak pattern. Suggests long signal awareness.
High-confidence prediction	AI prediction with 85%+ confidence.	Crisp banner validated by briefing with reference.
Account intelligence complete	All intelligence sources for [Account] are active.	Subtle banner indicating full intelligence coverage.

5.6 Action and CTA Copy Patterns

Action labels and CTA copy follow specific patterns. Actions are verbs that describe what the user will accomplish, not what the system will do. The user engages an account, the system does not. The user prepares for a meeting, the system does not. This distinction reinforces the "Intelligence as Executive Briefing" philosophy: the user is the decision-maker.

Action Type	Label Pattern	Examples
Primary action	Verb + Object	Engage Account, Prepare Briefing, View Evidence, Review Signals
Secondary action	Verb + Detail	See Full Reasoning, Expand Evidence, View Related Signals
Tertiary action	Object + Management	Save to Workspace, Dismiss Signal, Share Intelligence
Destructive action	Confirm + Consequence	Remove Source (Intelligence from this source will be removed), Clear History