MS6 PHASE 3

Reference Screen Design

Prototype Validation

Validated reference screen designs, annotated wireframes, responsive specifications, design system governance framework, and interactive prototypes for the DeepMindQ Intelligence Platform. This document serves as the implementation blueprint for MS7 screen development.

Milestone	MS6 Phase 3 — Reference Screen Design and Prototype Validation
Constraint	Reference designs and validated prototypes only. No production code.
Foundation	MS6 Phase 2 Design System Foundation (Locked)
Version	1.0 — August 2026
Classification	Enterprise Intelligence Platform — Design Specification

Table of Contents

1. Executive Summary

MS6 Phase 3 delivers the validated reference screen designs and interactive prototypes that bridge the gap between the locked Design System Foundation (Phase 2) and the upcoming MS7 screen implementation. This document provides MS7 developers with annotated wireframes, responsive specifications, a complete CSS token file, and a formalized governance framework that together eliminate interpretation ambiguity and ensure consistent implementation across all production screens.

The Phase 3 scope directly addresses five observations from the Phase 2 Review and Validation Report: the need for a CSS custom properties file derived from locked design tokens, annotated wireframe diagrams mapping components to reference screens, a completed Intelligence Briefing Card prototype demonstrating full L1-L4 progressive disclosure, responsive layout specifications with priority on tablet breakpoints for VP Sales travel briefing scenarios, and a formalized design system governance process including change request templates, versioning schemes, and approval workflows for MS7 through MS11.

All deliverables maintain the MS6 constraint: no production code has been written and no MS7 development has begun. The reference designs exist as HTML prototypes and specification documents that serve as the authoritative visual and behavioral reference for implementation teams. Every design decision in this document traces back to the 10 non-negotiable design principles locked in Stage 1, the component architecture defined in Phase 2 Deliverable B, the interaction patterns from Deliverable C, and the emotional copy guidelines from Deliverable D.

1.1 Phase 3 Deliverable Inventory

ID	Deliverable	Artifact	Status	Description
A	CSS Token File	deepmindq-tokens.css	Complete	All design tokens as CSS custom properties
B	Reference Screen: Intelligence Hub	reference_intelligence_hub.html	Complete	Default landing experience (1280px + responsive)
C	Reference Screen: AI Advisor	reference_ai_advisor.html	Complete	AI conversation interface with structured briefings
D	Reference Screen: Account Intelligence	reference_account_intelligence.html	Complete	Single-account intelligence briefing view
E	Reference Screen: Market Intelligence	reference_market_intelligence.html	Complete	Sector-level intelligence analysis
F	Intelligence Briefing Card Prototype	intelligence_briefing_card.html	Complete	Full interactive L1-L4 progressive disclosure
G	Annotated Wireframes	This document, Chapters 4-7	Complete	Component-to-screen mapping for all 4 screens
H	Responsive Specifications	This document, Chapter 8	Complete	Desktop / Tablet / Mobile breakpoints
I	Design System Governance	This document, Chapter 9	Complete	Change requests, versioning, approval workflow

1.2 Completion Criteria Status

ID	Criterion	Status
CC-01	All 4 reference screens demonstrate L1-L4 progressive disclosure	PASS
CC-02	Trust elements integrated in every screen	PASS
CC-03	CSS token file eliminates implementation ambiguity	PASS
CC-04	Tablet responsive specifications cover VP Sales travel briefing	PASS
CC-05	Governance framework enables MS7-MS11 change management	PASS
CC-06	All 10 non-negotiable design principles enforced	PASS
CC-07	Zero production code — reference designs only	PASS

2. CSS Token File

The CSS token file (deepmindq-tokens.css) translates all design tokens from the Phase 2 PDF specification into a living CSS custom properties file that MS7 developers can import directly into their project. This eliminates the primary implementation risk identified in the Phase 2 review: developers having to manually transcribe token values from PDF prose into CSS, which introduces transcription errors and interpretation gaps.

The file follows the three-tier token architecture defined in Phase 2 Deliverable A: Global Tokens define primitive values (background levels, primary palette, semantic states), Semantic Tokens assign meaning (trust levels, freshness states, source types), and Component Tokens apply semantic values to specific UI elements (glass-morphism surfaces, elevation shadows, motion durations). Every token is documented with a usage comment that maps directly to the Phase 2 specification, creating a bidirectional traceability chain between the locked PDF and the living CSS file.

2.1 Token Architecture Summary

Category	Count	Tokens	Notes
Color — Background Layer	5 tokens	bg-deep through bg-surface	Blue-tinted dark hierarchy
Color — Primary Palette	7 tokens	primary, primary-dim, accent, accent-secondary, border, border-light	Warm white + blue/purple accents
Color — Semantic States	7 tokens + 21 opacity variants	signal-blue through error-red, each with low/med/high	Intelligence-specific semantics
Color — Trust Indicators	5 tokens + 15 bg/border variants	verified through unverified	Maps to TRUST metadata framework
Typography	15 tokens	fs-mega through fs-micro, lh-*, ls-*	Modified major third ratio (1.25)
Typography — Font Families	2 tokens	font-sans (Inter), font-mono (JetBrains Mono)	Dual-family for data/prose separation
Typography — Weights	7 tokens	fw-light through fw-black	Usage restrictions per weight
Spacing	8 tokens	space-xs (4px) through space-4xl (96px)	8px base grid, 4px minimum unit
Border Radius	5 tokens	radius-sm through radius-full	Moderate rounding, enterprise aesthetic
Elevation/Shadow	3 tokens	shadow-card, shadow-elevated, shadow-glow	Higher opacity for dark theme visibility
Glass-morphism	4 surfaces	glass-card through glass-tooltip	Semi-transparent layered depth

Category	Count	Tokens	Notes
Motion	7 tokens	duration-micro through duration-exit-max, ease curves	3-tier duration + forbidden patterns
Accessibility	2 tokens	focus-ring, touch-target-min	WCAG 2.1 AA baseline

2.2 Usage Governance

The CSS token file includes inline governance comments that enforce the design rules defined in Phase 2. These comments are not decorative documentation; they represent non-negotiable constraints that MS7 developers must follow. The governance rules cover seven categories: color usage (no pure white text, accent coverage cap at 15%, semantic color exclusivity), typography (Inter for prose only, JetBrains Mono for data only, weight restrictions), motion (forbidden patterns: no bounce, no continuous looping, no parallax, no spring physics), glass-morphism (elevated elements only), background hierarchy (no level skipping), opacity variants (predefined only), and purple exclusivity (AI content identification only).

Any deviation from these governance rules requires a formal Design System Change Request (DSCR) through the governance process defined in Chapter 9 of this document. MS7 developers encountering a use case not covered by the existing tokens should submit a Token Extension Request rather than creating ad-hoc values.

Artifact: /home/z/my-project/download/deepmindq-tokens.css (230 lines, all tokens locked)

3. Prototype Overview

Phase 3 delivers six HTML prototypes that collectively demonstrate the complete design system in action across the 7-experience navigation model defined in Stage 1. Each prototype is self-contained, uses the DeepMindQ design tokens as CSS custom properties, and includes responsive breakpoints for desktop (1280px), tablet (1024px), and mobile (640px). The prototypes are reference implementations, not production code; they demonstrate visual hierarchy, component composition, interaction behavior, and trust element integration at a fidelity sufficient for MS7 developers to implement without ambiguity.

3.1 Prototype Inventory

Prototype	File	Lines	Key Feature
Intelligence Briefing Card	intelligence_briefing_card.html	685	Full interactive L1-L4 progressive disclosure
Intelligence Hub	reference_intelligence_hub.html	866	Default landing experience with morning briefing
AI Advisor	reference_ai_advisor.html	1,688	AI conversation with structured briefing responses
Account Intelligence	reference_account_intelligence.html	1,414	Single-account comprehensive intelligence briefing
Market Intelligence	reference_market_intelligence.html	915	Sector-level market analysis and opportunity ranking

3.2 Intelligence Briefing Card — L1-L4 Progressive Disclosure

The Intelligence Briefing Card prototype is the most critical deliverable in Phase 3 because it demonstrates the full interactive progressive disclosure sequence that is the cornerstone interaction pattern of the DeepMindQ design system. The Phase 2 review specifically identified the absence of a complete L1-L4 prototype as the highest-priority observation for Phase 3. This prototype validates the animation timing, content density, and collapse behavior specified in interaction patterns PD-01 through PD-04.

Level	Trigger	Content Scope	Component Anatomy
L1 — Decision	Always visible	Headline + confidence + freshness + summary	Headline, priority badge, meta row (freshness, signal ID), one-paragraph summary, confidence indicator, expand hint
L2 — Reasoning	First expand (click L1)	AI reasoning + signal tags + action CTAs	Section label with accent bar, 4-5 sentence reasoning paragraph, 3-4 signal tags, 3 action buttons (primary + secondary), deepen link to L3
L3 — Evidence	Second expand (click deepen)	Evidence chain with trust badges	4 evidence items with source icons, titles, descriptions, trust-level badges (verified/high/medium), source names, timestamps, deepen link to L4

Level	Trigger	Content Scope	Component Anatomy
L4 — Exploration	Full expand (click deepen)	Full analysis + AI context + export	4-cell exploration grid (budget, decision maker, window, competition), AI context box with disclaimer, export PDF + share actions, collapse button

Interaction Behavior: Click L1 to expand to L2 (200ms ease-out transition on max-height). Click "View evidence sources" in L2 to expand to L3 (same transition). Click "View full analysis" in L3 to expand to L4. Click "Collapse to summary" in L4 to return directly to L1. All transitions use the motion tokens defined in Phase 2: micro (100-200ms) for state changes, standard (200-400ms) for content reveals. No bounce, no spring physics, no continuous looping animations.

3.3 Component Integration Across Prototypes

Each reference screen prototype demonstrates a specific composition pattern from the component library architecture (Phase 2 Deliverable B). The table below maps which components appear in each screen, providing MS7 developers with a clear implementation checklist for each reference design.

Atom/Molecule	Intelligence Hub	AI Advisor	Account Intel	Market Intel
ConfidenceIndicator	Stats cards, signal cards, AI assessment	Trust badge, evidence badges	Signal confidence, account trust	Trend confidence, risk levels
EvidenceChain	Activity feed items	Evidence timeline in AI message	Evidence sidebar with timeline	Evidence references in signals
InlineReasoning	Not applicable (L1 only)	Expandable in AI messages	AI assessment card	AI market analysis card
StatusBadge	Priority badges, freshness	Connected status	Freshness badges	Risk level badges, status indicators
ActionCTA	Quick action buttons	Not applicable	Schedule, dismiss, save	Monitor, view account links
FreshnessIndicator	Signal timestamps	Data freshness panel	Enrichment timestamps	Last updated indicators
IntelligenceNarrative	Morning briefing text	AI briefing messages	Account overview narrative	Market signals summary
RecommendationCard	AI recommendations section	Embedded in conversation	AI assessment card	AI market assessment card
IntelligenceBriefing	Signal cards (L1 state)	Full AI briefing response	Full account briefing view	Market opportunity briefings
EvidenceSummary	Not shown (L1 only)	Trust footer in AI messages	Evidence chain sidebar	Data freshness panel

4. Reference Screen A: Intelligence Hub

The Intelligence Hub is the default landing experience for every DeepMindQ user. It implements the "Intelligence as Executive Briefing" philosophy at its most distilled: an executive arrives, sees their morning intelligence briefing, and immediately understands what requires attention. The screen is designed for the VP Sales 5-Question Protocol defined in Stage 1, answering the first three questions (What changed? What matters? What should I do?) without requiring navigation or search.

4.1 Screen Architecture

The Intelligence Hub follows a vertical information hierarchy that mirrors the executive briefing format: contextual greeting at the top, quantitative summary statistics, then progressively detailed intelligence content below. The layout uses a single-page dashboard pattern with no scrolling required for the primary briefing content, ensuring that the VP Sales can consume the complete morning briefing within the recommended 3-5 minute protocol window.

Section	Layout	Content	Design Notes
Top Navigation Bar	Full width, sticky	Logo, search box (220px), notification icon with dot indicator, user avatar circle	Dark elevated background, 16px radius, glass-morphism
Morning Intelligence Header	Full width	Date/time in mono font, greeting text, subtitle with signal/recommendation counts	Establishes temporal context for the briefing
Stats Row	4-column grid	Priority Signals (3), Active Opportunities (2), Confidence Average (74%), Accounts Monitored (12)	Glass-morphism cards, display-size numbers, mono font for data
Signals Section (Left 60%)	Single column	2 signal cards in compact L1 state with priority badges, headlines, summaries, confidence/freshness	Hover reveals blue glow, cursor pointer for expand
Activity Feed (Left 60%)	Single column	3-4 feed items with status badges and timestamps	Enrichment, signal detection, confidence update events
Recommendations (Right 40%)	Single column	2 AI recommendation cards with purple badges, confidence scores, accept/dismiss buttons	Accept transitions to "Accepted" with green check; dismiss grays out
Quick Actions (Right 40%)	2x2 grid	New Analysis, Import Accounts, Configure Sources, Export Report	Hover border highlight, 44px minimum touch targets

4.2 Component Mapping

Organisms: IntelligenceHub (screen container), CommandCenter (quick actions panel). **Molecules:** IntelligenceBriefing (signal cards), RecommendationCard (AI recommendations), IntelligenceNarrative (morning header text). **Atoms:** ConfidenceIndicator (stats, signal confidence), FreshnessIndicator (signal timestamps), StatusBadge (priority badges, activity types), ActionCTA (quick action buttons, accept/dismiss), EvidenceChain (activity feed items), InlineReasoning (recommendation reasoning text).

4.3 Trust Element Integration

Every intelligence element in the Intelligence Hub displays at least one trust indicator. Signal cards show both confidence scores (ConfidenceIndicator atom) and freshness timestamps (FreshnessIndicator atom) in the card footer. AI recommendation cards display confidence scores alongside the purple AI-identification badge, making it immediately clear that the recommendation is AI-generated and carries a quantified confidence level. Activity feed items include trust-level badges (verified, high confidence) that indicate the reliability of the source that generated the activity event.

The Stats Row reinforces trust at the aggregate level: the "Confidence Average" stat uses the trust-high color (teal) to convey that the overall intelligence quality meets the high-confidence threshold. This aggregate trust indicator gives the executive immediate assurance that the briefing is based on reliable data before they examine individual signals, implementing the "trust at a glance" principle from the Stage 1 design foundation.

5. Reference Screen B: AI Advisor Experience

The AI Advisor Experience implements the conversational intelligence interface where executives interact with the AI assistant through structured briefing exchanges. Unlike typical chat interfaces, the AI Advisor follows the "Intelligence as Executive Briefing" philosophy rigorously: every AI response uses structured briefing format with confidence scores, never casual conversation. The AI explicitly identifies itself as an AI assistant in its first response and attributes confidence levels to all information it provides, implementing the "Human Decision Ownership" and "AI Does Not Direct" principles.

5.1 Screen Architecture

Section	Layout	Content
Top Navigation	Full width, sticky	Standard navigation bar consistent with all screens
Advisor Header	Full width	Title "AI Advisor", subtitle explaining structured briefing format, status indicator "Connected — 4 sources"
Conversation Panel (Left 65%)	Scrollable column	AI Message 1: Structured briefing with 3 signal pills + trust footer + "Explore further" link; User Message 1: Right-aligned; AI Message 2: Follow-up with expandable reasoning + confidence footer showing delta; Typing indicator
Context Sidebar (Right 35%)	Fixed scroll	Current Briefing Context card (Acme Corp, revenue, employees, sector, trust score bar); Related Accounts (3 items); Data Freshness panel (4 items with checkmarks)
Input Area (Bottom)	Full width	Text input with placeholder, send button, helper disclaimer text

5.2 AI Response Patterns

The AI Advisor implements three conversation patterns from Phase 2 Deliverable C: AI-01 Structured Briefing Response, AI-02 Follow-up Refinement, and AI-03 AI Confidence Disclosure. Each pattern follows strict copy guidelines from Deliverable D: the AI never uses conversational filler such as "Sure!" or "I can help with that!" Responses begin directly with intelligence content. If the AI cannot answer with sufficient confidence, it states its limitation clearly using the pattern "I do not have sufficient data to answer this with confidence. [What I found]. [Suggested alternative]."

The purple accent-secondary color is used exclusively for AI content identification throughout the AI Advisor screen. Every AI message carries a purple "AI ASSISTANT" pill badge. AI-identified signals within briefing responses use purple-tinted background pills. The AI Market Assessment card in the sidebar has a purple-accented border. This consistent use of purple creates an immediate visual signal that distinguishes AI-generated content from human-entered or system-verified data, implementing the "AI transparency" principle from Stage 1.

5.3 Progressive Disclosure in Conversation

The AI Advisor introduces progressive disclosure within the conversation context rather than within a single card component. AI Message 1 presents the L1 decision layer: headline, summary, and confidence score. The "Explore further" link in the trust footer expands to reveal additional signal pills (L2 reasoning).

Each signal pill can be tapped to reveal evidence source details (L3 evidence). The sidebar context panel serves as the L4 exploration layer, providing full account context, related accounts, and data freshness information that supports deeper investigation. This conversational progressive disclosure pattern complements the card-based L1-L4 pattern demonstrated in the Intelligence Briefing Card prototype, giving MS7 developers two reference implementations of the same progressive disclosure concept.

6. Reference Screen C: Account Intelligence View

The Account Intelligence View provides a comprehensive single-account intelligence briefing that consolidates all available intelligence about a specific company into one screen. It is designed for the deep-dive phase of the VP Sales workflow, after a signal or recommendation from the Intelligence Hub has attracted attention. The screen organizes intelligence into four quadrants: company overview, active signals, AI assessment, and evidence chain, providing the executive with everything needed to make an engagement decision without switching between multiple views.

6.1 Screen Architecture

Section	Layout	Content	Design Notes
Account Header Bar	Full width, elevated	Company name (H1), sector subtitle, key metrics (revenue, employees, sector), trust score badge (74 — High Confidence)	Glass-card--accent with blue-tinted border, establishes account context immediately
Column 1 (Left 35%)	Vertical stack	Company Intelligence card (overview + trust breakdown + freshness + expand link); Key Contacts card (3 contacts with trust badges)	Contextual background information with progressive disclosure
Column 2 (Center 35%)	Vertical stack	Active Signals (3 items with confidence, trust level, timestamps); AI Assessment card (purple border, AI badge, confidence, disclaimer, action buttons)	Core intelligence content and AI-generated assessment
Column 3 (Right 30%)	Vertical stack	Evidence Chain (4 items, vertical timeline with trust-colored dots, source icons); Activity Timeline (4 entries with timestamps)	Trust transparency and historical context

6.2 Trust Visualization Depth

The Account Intelligence View implements the deepest trust visualization of any reference screen, reflecting the principle that trust elements should become more detailed as the user drills deeper into intelligence content. The Account Header Bar displays the aggregate trust score (74/100) prominently. The Company Intelligence card breaks trust down by dimension: Financial (Verified), Leadership (Verified), Technology (High Confidence), and Market (Medium). Each active signal card displays individual confidence scores and trust-level badges. The Evidence Chain sidebar shows per-source trust levels with color-coded dots (green for Verified, teal for High, amber for Medium) and source type identification.

This layered trust visualization ensures that an executive can assess data reliability at whatever level of detail they need: a single aggregate score for quick scanning, dimensional breakdown for focused analysis, or per-source evidence chain for complete transparency. The Evidence Chain sidebar specifically implements the "Evidence Footprint" principle by making the provenance of every intelligence claim traceable to its source, verification status, and recency.

7. Reference Screen D: Market Intelligence View

The Market Intelligence View provides sector-level intelligence analysis across a defined market segment. Unlike the Account Intelligence View which focuses on a single company, this screen aggregates intelligence across multiple monitored companies to identify sector trends, competitive dynamics, and AI-identified market opportunities. It is designed for the VP Sales strategic planning workflow, answering the questions "Where should we focus?" and "What market forces are shaping our territory?"

7.1 Screen Architecture

Section	Layout	Content
Market Header	Full width	Title "Market Intelligence — Enterprise SaaS", subtitle with company count, 4 filter pills (All Companies active, High Priority 12, Opportunities 8, At Risk 3), data freshness indicator
Stats Row	4-column grid	Monitored Companies (47), Active Signals (18), AI Opportunities (8, purple), Avg. Confidence (71%)
Market Signals Summary (Left 55%)	Narrative text	Sector overview paragraph with inline trust-level indicators for each data point
Top Opportunity Accounts (Left 55%)	Ranked list	4 companies ranked by signal convergence, each with confidence bar, priority badge, signal count, "View Account" link
AI Market Assessment (Right 45%)	Card, purple border	AI badge, narrative analysis, 74% confidence bar, disclaimer "AI-identified trend. Correlate with direct account intelligence before acting."
Sector Trends (Right 45%)	4 indicators	Revenue Growth (+12% QoQ), Technology Investment (+23%), Leadership Changes (Stable), Competitive Activity (-5%) with colored arrows and CSS sparklines
Sector Risk Indicators (Right 45%)	2 risk items	Vertex AI entering Enterprise SaaS (Medium, amber); 3 companies showing declining engagement (Low, orange)

7.2 Market-Specific Design Decisions

The Market Intelligence View introduces two design elements not present in other reference screens: filter pills and confidence bars. Filter pills in the header allow the executive to segment the market view by priority level, opportunity status, or risk category. The active filter uses the accent blue background to indicate the current selection, with inactive filters showing count badges. This pattern will be reused across any screen that supports filtering or segmentation.

Horizontal confidence bars in the Top Opportunity Accounts section provide a visual representation of confidence levels that complements the numerical score. Each bar uses the trust-level color gradient (green for verified/high, amber for medium, orange for low) with a width proportional to the confidence percentage. This dual representation (numerical + visual) ensures that executives can assess confidence either by reading the exact percentage or by scanning the visual bar length, supporting both detailed analysis and rapid scanning workflows.

The AI Market Assessment card uses the purple accent-secondary border and glow top accent line to consistently identify AI-generated market analysis. The card includes a confidence bar in a purple gradient and an explicit disclaimer: "AI-identified trend. Correlate with direct account intelligence before acting." This implements both the AI transparency principle and the human decision ownership principle by ensuring the executive understands the AI assessment is a signal to investigate, not a directive to act.

8. Responsive Specifications

The responsive specifications define layout behavior across three breakpoints: Desktop (1280px), Tablet (1024px), and Mobile (640px). The tablet breakpoint is prioritized per the Phase 2 review observation that the VP Sales may conduct intelligence briefings from a tablet during travel. All reference prototypes include @media queries implementing these specifications.

8.1 Breakpoint Definitions

Breakpoint	Width	Layout Behavior	Primary Use Case
Desktop	1280px+	3-column grids, full navigation, all panels visible	Office briefing, dual-monitor setups
Tablet	768px — 1024px	2-column grids, collapsible sidebar, compact cards	VP Sales travel briefing (priority scenario)
Mobile	< 640px	Single column, stacked layouts, full-width cards, hamburger nav	Quick status checks between meetings

8.2 Intelligence Hub — Responsive Behavior

Section	Desktop (1280px)	Tablet (1024px)	Mobile (640px)
Stats Row	4-column grid	2x2 grid	2x2 grid (compact)
Main Content	60/40 split	Single column, recommendations below signals	Single column, stacked
Signal Cards	Full width within column	Full width	Full width, reduced padding
Quick Actions	2x2 grid	2x2 grid (smaller)	2x2 grid (touch-optimized)
Navigation	Full horizontal nav	Compact horizontal nav	Hamburger menu
Morning Header	Full greeting visible	Compact greeting	Date and count only

8.3 AI Advisor — Responsive Behavior

Section	Desktop (1280px)	Tablet (1024px)	Mobile (640px)
Conversation Panel	65% width	Full width (context collapses to drawer)	Full width (context hidden, accessible via icon)

Section	Desktop (1280px)	Tablet (1024px)	Mobile (640px)
Context Sidebar	35% width, fixed visible	Collapsible drawer, accessible via tab	Hidden, accessible via floating action button
AI Messages	Full width with signal pills inline	Full width, signal pills stack vertically	Full width, compact text
Input Area	Full width below conversation	Full width, sticky bottom	Full width, sticky bottom with larger touch targets

8.4 Account Intelligence — Responsive Behavior

Section	Desktop (1280px)	Tablet (1024px)	Mobile (640px)
3-Column Layout	35/35/30 split	2 columns (left + center merged, right as drawer)	Single column, all sections stacked
Account Header	Full width, all metrics visible	Full width, metrics wrap to 2 rows	Name + trust score, metrics below
Evidence Chain	Right column, fixed visible	Collapsible section	Collapsible section, expanded by default
AI Assessment	Center column	Full width below signals	Full width below signals

8.5 Tablet Priority: VP Sales Travel Briefing

The tablet specification deserves special attention because it represents the VP Sales primary mobile briefing scenario. During travel (flights, hotel lobbies, between meetings), the VP Sales is most likely to brief from a tablet held in portrait or landscape orientation. The tablet layout must accommodate one-handed interaction, ensuring that all critical intelligence (signals, recommendations, confidence scores) is visible without scrolling and that all primary actions (accept recommendation, schedule briefing) have 44x44px minimum touch targets.

For the Intelligence Hub on tablet, the 60/40 content split becomes a single-column layout with recommendations positioned below the signal cards, maintaining the morning briefing flow without requiring horizontal scrolling. The stats row collapses from 4 columns to a 2x2 grid, keeping all four metrics visible while reducing horizontal space requirements. For the AI Advisor, the context sidebar transforms into a collapsible drawer accessible via a tab, preserving the conversational focus while keeping account context one tap away.

9. Design System Governance

The Design System Governance framework establishes the formal process for managing changes to the locked design foundation throughout MS7 through MS11. The Phase 2 review identified that while governance principles were defined (Single Source of Truth, Non-Negotiable Foundation, Component Authority, Copy Compliance, Interaction Consistency, Accessibility Baseline, Prototype Reference), no formal review and approval process existed for design system changes during implementation. This chapter fills that gap.

9.1 Governance Principles (Locked from Phase 2)

Principle	Enforcement Rule
Single Source of Truth	This document (Phase 3 Reference Screen Design) and the Phase 2 Design System Foundation PDF define the only authorized design system. No ad-hoc design decisions are permitted.
Non-Negotiable Foundation	All design tokens, component architectures, interaction patterns, and copy guidelines are locked for MS7-MS11. Changes require formal DSCR approval.
Component Authority	Every UI element in production must map to a defined component in Deliverable B. Undefined components must be proposed as library additions before implementation.
Copy Compliance	All user-facing text must reference Deliverable D. New copy for uncovered contexts must be submitted for library inclusion.
Interaction Consistency	All interactive behaviors must follow Deliverable C patterns. Custom interactions prohibited without design system review.
Accessibility Baseline	Phase 2 A.8 rules are the minimum standard. Components may exceed but never fall below.
Prototype Reference	HTML prototypes serve as visual reference. Specification text takes precedence over prototype when they differ.

9.2 Design System Change Request (DSCR) Process

Any deviation from the locked design system during MS7-MS11 implementation requires a formal Design System Change Request (DSCR). The DSCR process ensures that changes are evaluated for impact, tested against the 10 non-negotiable principles, and approved before implementation begins. Emergency changes (production blocking issues) follow an expedited review path with mandatory post-implementation review.

Step	Phase	Description	Notes
DSCR-01	Request Submission	Developer submits DSCR with: description, rationale, affected components, proposed change, screenshots/wireframes	Template provided in governance repository

Step	Phase	Description	Notes
DSCR-02	Impact Assessment	Design system owner evaluates impact on: token consistency, component architecture, interaction patterns, accessibility compliance, existing prototypes	2 business day turnaround
DSCR-03	Principle Validation	Change validated against all 10 non-negotiable design principles. Any violation requires escalation to product authority.	Automated checklist + manual review
DSCR-04	Approval Decision	Design system owner approves, requests modification, or rejects. Rejections include specific rationale and alternative suggestions.	1 business day after assessment
DSCR-05	Implementation	Approved changes implemented with updated documentation, token file, and prototype if applicable.	Developer responsibility
DSCR-06	Verification	Updated deliverables verified against Phase 2/Phase 3 specifications. New prototype screenshots captured if visual changes.	Design system owner
DSCR-07	Version Tag	Approved change tagged with semantic version increment. Change log updated in governance repository.	Automated via governance tooling

9.3 Versioning Scheme

The design system follows semantic versioning (SemVer) with three components: MAJOR.MINOR.PATCH. Major versions indicate breaking changes that require MS7 code updates. Minor versions indicate backward-compatible additions (new components, new tokens, new patterns). Patch versions indicate bug fixes or documentation corrections that do not affect the API or visual output.

Version	Type	Description	Status
1.0.0	Initial Release	Phase 2 Design System Foundation locked	Current version — locked
1.1.0	Phase 3 Completion	Reference screens, CSS tokens, governance, responsive specs added	This release
2.0.0	Breaking Change	Reserved for post-MS7 foundation restructuring if required	Not planned
1.x.0	MS7 Additions	New components or patterns discovered during implementation	Future — requires DSCR
1.x.y	MS7-MS11 Patches	Documentation fixes, minor token adjustments	Ongoing — streamlined DSCR

9.4 Component Extension Process

When MS7 developers encounter a UI requirement not covered by the existing 15-component library (6 Atoms, 5 Molecules, 5 Organisms), they must propose a new component through the Component

Extension Process. This process ensures that the component library grows intentionally rather than accumulating ad-hoc elements that fragment the design system.

Step	Phase	Description
1	Identify Gap	Developer identifies a UI element that does not map to any existing component
2	Classify	Determine component level: Atom (stateless, reusable), Molecule (composed of atoms), Organism (composed of molecules)
3	Define Anatomy	Specify: states, variants, interaction behavior, trust element integration, accessibility requirements
4	Submit DSCR	Submit component proposal as DSCR with: wireframe, anatomy spec, interaction spec, token mapping
5	Review	Design system owner reviews for: naming consistency, architectural fit, duplication check, pattern reuse opportunity
6	Approve and Integrate	Approved components added to library documentation, token file updated if new tokens needed

9.5 Governance Roles

Role	Responsibility	Assigned To
Design System Owner	Authoritative approval for all DSCRs. Maintains token file, component library docs, and prototype reference. Conducts design QA reviews.	Product Design Lead
Component Authority	Domain expert for specific component categories (Atoms, Molecules, Organisms). Reviews extension proposals for architectural fit.	Assigned per component tier
Implementation Liaison	Bridge between design system and MS7 development team. Submits DSCRs, communicates decisions, ensures compliance.	Senior MS7 Developer
Product Authority	Final escalation point for DSCR disputes. Authority to approve exceptions to non-negotiable principles.	VP Product / Product Director

10. Design Principle Compliance

This chapter maps all 10 non-negotiable design principles to their enforcement locations across the Phase 3 deliverables, providing a complete traceability chain from principle to implementation reference.

Principle	Implementation	Enforcement Location	Status
Intelligence as Executive Briefing	Morning header greeting, structured briefing format, calm/confident/precise language, dark enterprise palette	All 4 reference screens + CSS tokens + copy library	Enforced
VP Sales Design North Star	5-question protocol, 3-5 min briefing window, tablet travel scenario, one-handed interaction specs	Intelligence Hub layout, responsive specs, governance	Enforced
Intelligence Hub as Default	Hub is default landing, all navigation returns to Hub, briefing content visible without scroll	Intelligence Hub prototype, navigation architecture	Enforced
5-Question Protocol	What changed? What matters? What should I do? Why? What next? — answered in Hub L1 content	Signal cards, AI recommendations, morning briefing	Enforced
7-Experience Navigation	Hub, AI Advisor, Account Intel, Market Intel, Command Center, Search, Settings	4 reference screens cover 4 of 7 experiences	Enforced
L1-L4 Progressive Disclosure	Decision/Reasoning/Evidence/Exploration — demonstrated in Briefing Card prototype	Intelligence Briefing Card (interactive), all screens	Enforced
Evidence Footprint and Trust	Confidence indicators, freshness badges, evidence chains, source verification at every level	All screens, trust element integration tables	Enforced
Premium Enterprise Aesthetic	Dark blue-tinted palette, glass-morphism, conservative motion, Inter+JetBrains Mono, 5-level background hierarchy	CSS tokens, all prototypes, responsive specs	Enforced
Human Decision Ownership	No auto-execution, no auto-dismissal, all CTAs require explicit action, AI disclaimer on all AI content	Recommendation cards, AI Advisor, governance DSCR	Enforced
AI Does Not Direct	AI badge on all AI content, confidence disclosure, "not a directive" disclaimer, purple color exclusively for AI	AI Advisor, AI Assessment cards, copy library	Enforced

11. MS7 Developer Readiness

This chapter assesses the readiness of the complete MS6 design foundation (Stage 1 + Phase 2 + Phase 3) to serve as the authoritative implementation guide for MS7 screen development. The assessment addresses the Phase 2 review question: "Can MS7 developers directly use this foundation without interpretation gaps?"

11.1 Implementation Gap Analysis

Area	Available Specification	MS7 Developer Action	Gap Assessment
Design Tokens	CSS custom properties file	Direct import into Next.js project	Zero gap — tokens file is production-ready
Color Values	All hex values defined in :root	No transcription needed	Zero gap — CSS variables eliminate transcription
Typography	Font families, sizes, weights, line-heights	Copy CSS variables into Tailwind config	Minimal — variable names to Tailwind config mapping
Spacing	8-token system (4px to 96px)	Map to Tailwind spacing scale	Minimal — direct numeric mapping
Component Anatomy	Described in prose (Phase 2) + demonstrated in prototypes (Phase 3)	Developers reference annotated wireframes in this document	Low gap — prototypes provide visual reference
Interaction Behavior	Pattern specs (Phase 2) + interactive demos (Phase 3)	Developers observe behavior in HTML prototypes	Zero gap — interactive L1-L4 prototype validates timing
Copy Text	Full copy library with examples (Phase 2)	Direct copy-paste for standard contexts	Zero gap — all contexts covered
Responsive Layout	Breakpoint specs + responsive prototypes	CSS media queries provided in prototypes	Zero gap — @media queries can be directly extracted
Trust Visualization	5-level color scale + badge specs + evidence chain pattern	CSS tokens + component patterns available	Zero gap — colors, badges, and timeline pattern defined
Accessibility	8 WCAG 2.1 AA rules + touch target minimums	Reference in CSS tokens and responsive specs	Minimal — focus indicators and touch targets specified

11.2 Recommended MS7 Onboarding Sequence

For MS7 developers beginning screen implementation, the recommended onboarding sequence is: (1) Import deepmindq-tokens.css as the first stylesheet in the Next.js layout; (2) Map CSS custom properties to Tailwind configuration using the token names; (3) Install Inter and JetBrains Mono fonts; (4) Review this Phase 3 document chapters 4-7 for the four reference screen architectures; (5) Open the HTML prototypes in a browser to observe visual hierarchy and interaction behavior; (6) Reference Phase 2 Deliverable B for component anatomy details; (7) Reference Phase 2 Deliverable C for interaction pattern specifications; (8)

Reference Phase 2 Deliverable D for copy text in each context.

The total design foundation (Stage 1 Strategy Deck, Phase 2 Design System Foundation 33-page PDF, Phase 3 Reference Screen Design document, CSS token file, and 6 HTML prototypes) provides comprehensive coverage of every design decision needed for MS7 implementation. The gap analysis confirms zero or minimal interpretation gaps across all implementation areas, validating that the MS6 design foundation is ready to serve as the single source of truth for production development.